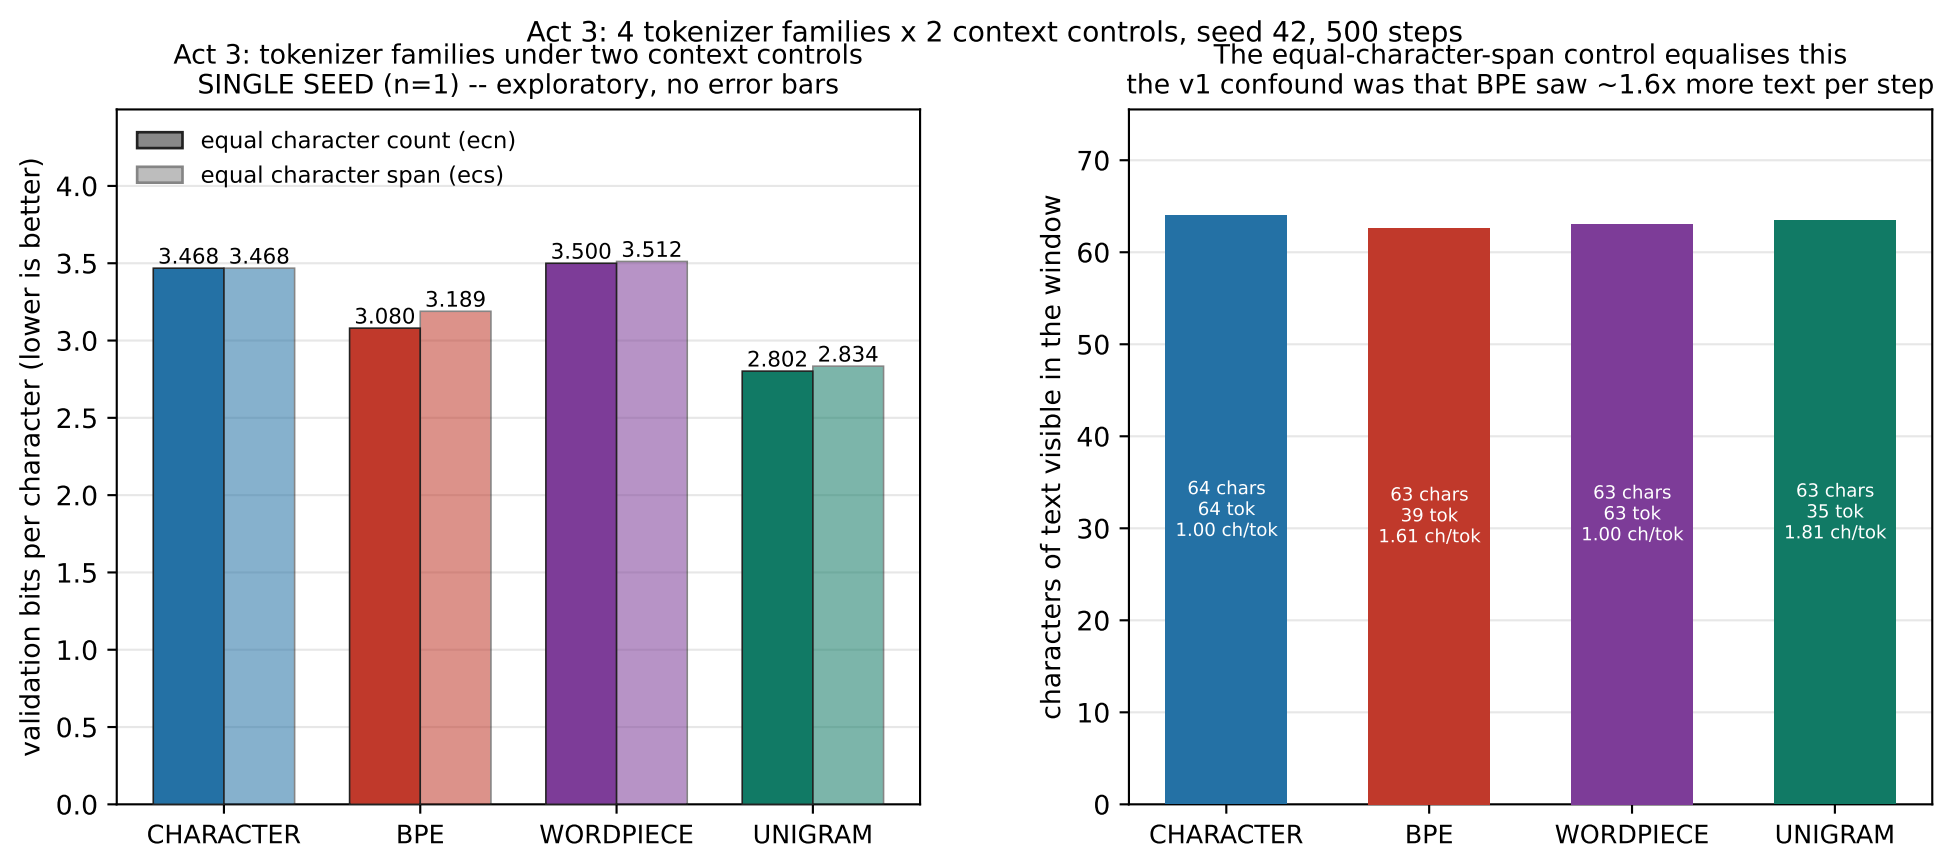
Act 3: 4 tokenizer families x 2 context controls, seed 42, 500 steps
Act 3: tokenizer families under two context controls
SINGLE SEED (n=1) -- exploratory, no error bars
equal character count (ecn)
equal character span (ecs)
validation bits per character (lower is better)
0.0
0.5
1.0
1.5
2.0
2.5
3.0
3.5
4.0
3.468
3.468
3.080
3.189
3.500
3.512
2.802
2.834
CHARACTER
BPE
WORDPIECE
UNIGRAM
The equal-character-span control equalises this
the v1 confound was that BPE saw ~1.6x more text per step
characters of text visible in the window
0
10
20
30
40
50
60
70
64 chars
64 tok
1.00 ch/tok
63 chars
39 tok
1.61 ch/tok
63 chars
63 tok
1.00 ch/tok
63 chars
35 tok
1.81 ch/tok
CHARACTER
BPE
WORDPIECE
UNIGRAM