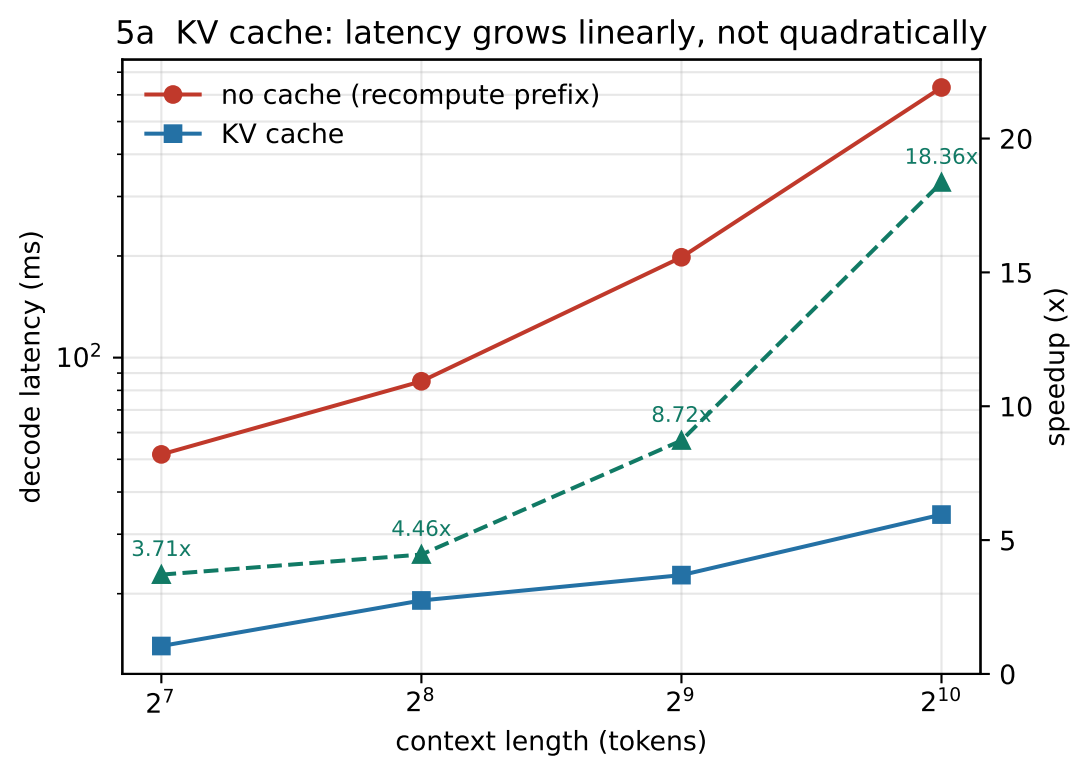
5a KV cache: latency grows linearly, not quadratically
no cache (recompute prefix)
KV cache
18.36x
8.72x
4.46x
3.71x
decode latency (ms)
speedup (x)
context length (tokens)

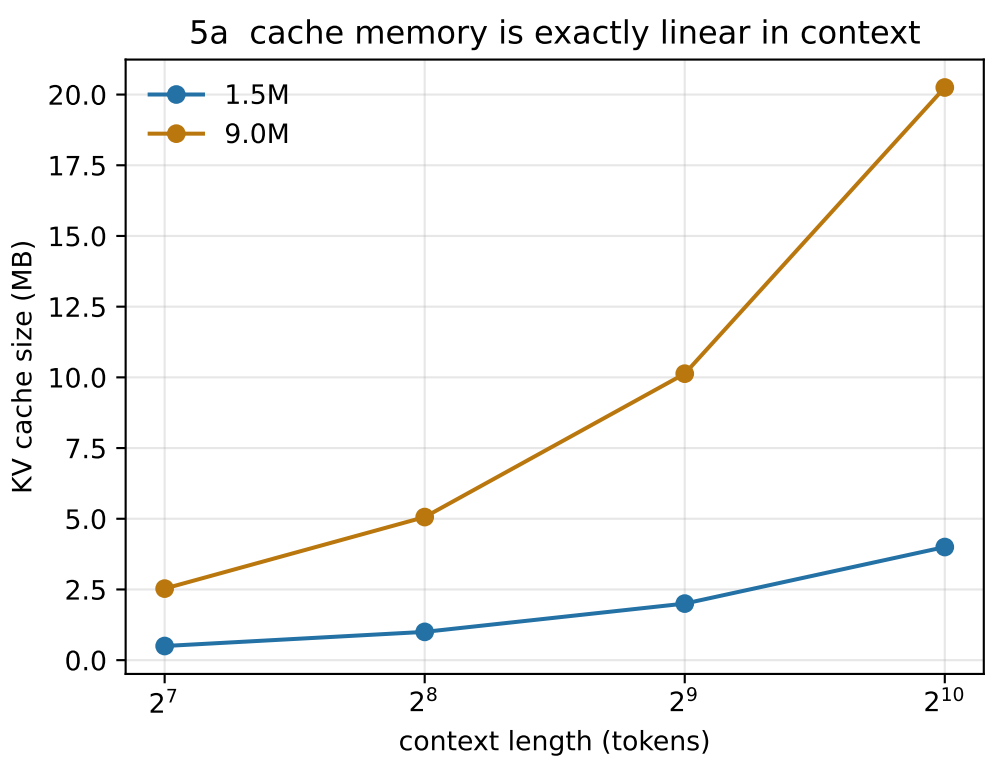
5a cache memory is exactly linear in context
1.5M
9.0M
KV cache size (MB)
context length (tokens)