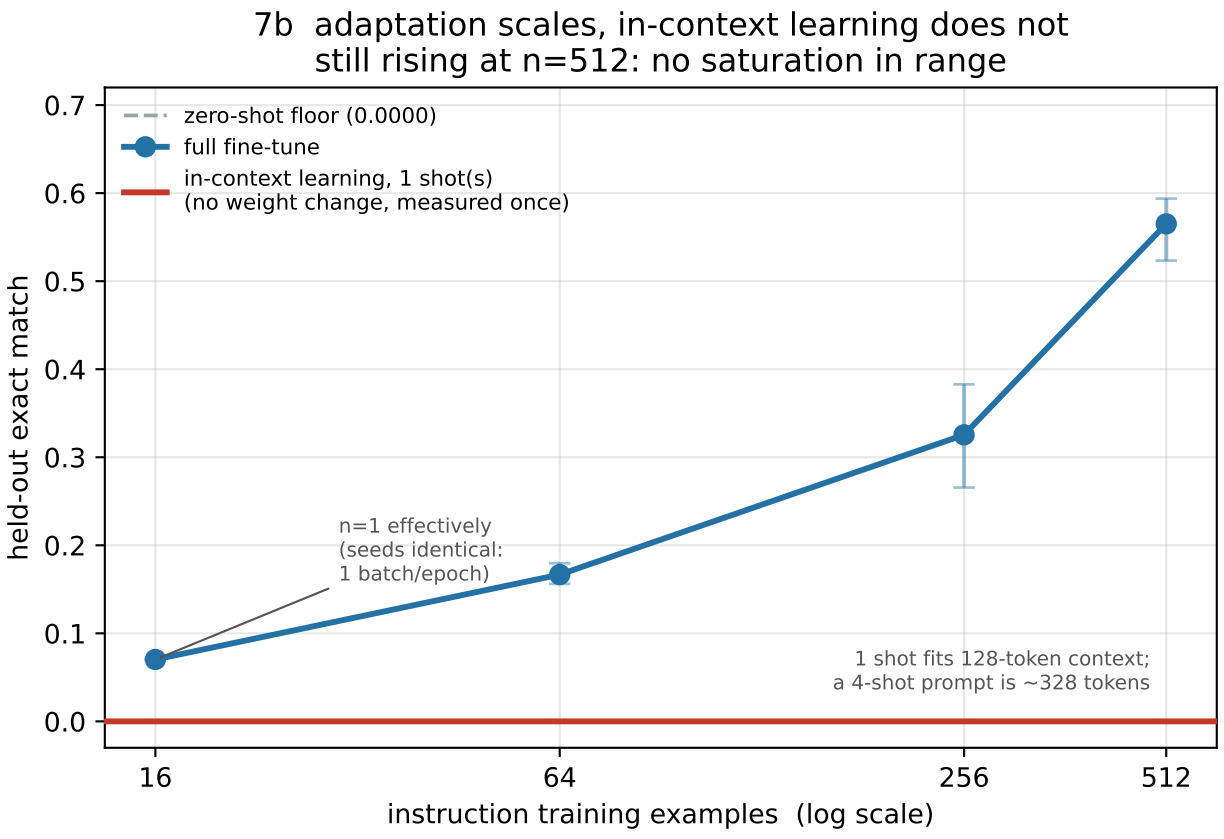
7b adaptation scales, in-context learning does not
still rising at n=512: no saturation in range
zero-shot floor (0.0000)
full fine-tune
in-context learning, 1 shot(s)
(no weight change, measured once)
n=1 effectively
(seeds identical:
1 batch/epoch)
1 shot fits 128-token context;
a 4-shot prompt is ~328 tokens
held-out exact match
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
16
64
256
512
instruction training examples (log scale)